

CLOUD COMPUTING PROJECT

CUSTOMER DATA DATABASE MIGRATION FROM ON-PREMISE TO CLOUD

OBJECTIVES:

After performing the project, the user should be able to:

1. Design a platform for recording, organizing, and analytics on customers transactions data received through application(s).
2. The data stored in the DB must be organized and maintained in such a way that caters to the needs of the future. Derive insights from the results to device better strategies and enhance customers experience.

THEORY:


You can use AWS Database Migration Service (AWS DMS) to migrate your data to and from most widely used commercial and open-source databases such as Oracle, PostgreSQL, Microsoft SQL Server, Amazon Redshift, Amazon Aurora, MariaDB, and MySQL. The service supports homogeneous migrations such as Oracle to Oracle, and also heterogeneous migrations between different database platforms, such as Oracle to MySQL or MySQL to Amazon Aurora with MySQL compatibility. The source or target database must be on an AWS service.

This pattern provides guidance for migrating an on-premises MySQL database to a MySQL database on an Amazon RDS. The pattern discusses the use of AWS Database Migration Service tools such as mysqldbcopy, mysqldump and mysql workbench migration methodology.

SOFTWARE NEEDED

MySQL Workbench
AWS RDS MySQL Environment
Data to be used

MIGRATION DESIGN AND SCHEMA DESIGN



tblcustomer	
order_id	VARCHAR(12)
Name	VARCHAR(12)
email_id_	VARCHAR(25)
Date	DATE
Fruits	VARCHAR(4000)
Vegetables	VARCHAR(4000)
Milk	VARCHAR(4000)
Rice	VARCHAR(4000)
Basket	TEXT
Indexes	

1. Steps in Importing file from .csv file to MySQL On-Premise

1. Start Excel, select the **Data** menu tab, and then click **MySQL for Excel** to open the MySQL for Excel task pane.
2. From the **Open a MySQL Connection** area in the task pane, double-click an existing local or remote connection to display the available database schemas.
3. Select a schema from the list and click **Next** to display all database objects in the schema (tables, views, and procedures).
4. Select the table, view, or procedure with data to import and then click **Import MySQL Data**. A preview window displays the selected data and provides **Options** and Advanced Options to be used during the import operation.
5. Click **Import** to finish the operation.

The Import Data windows provides a preview of the columns to select during the import operation. You can specify both the columns and rows to import. As the following figure shows, the preview includes a small subset of the rows for the selected table or view.

Figure 1.1 Importing table data with MySQL for Excel

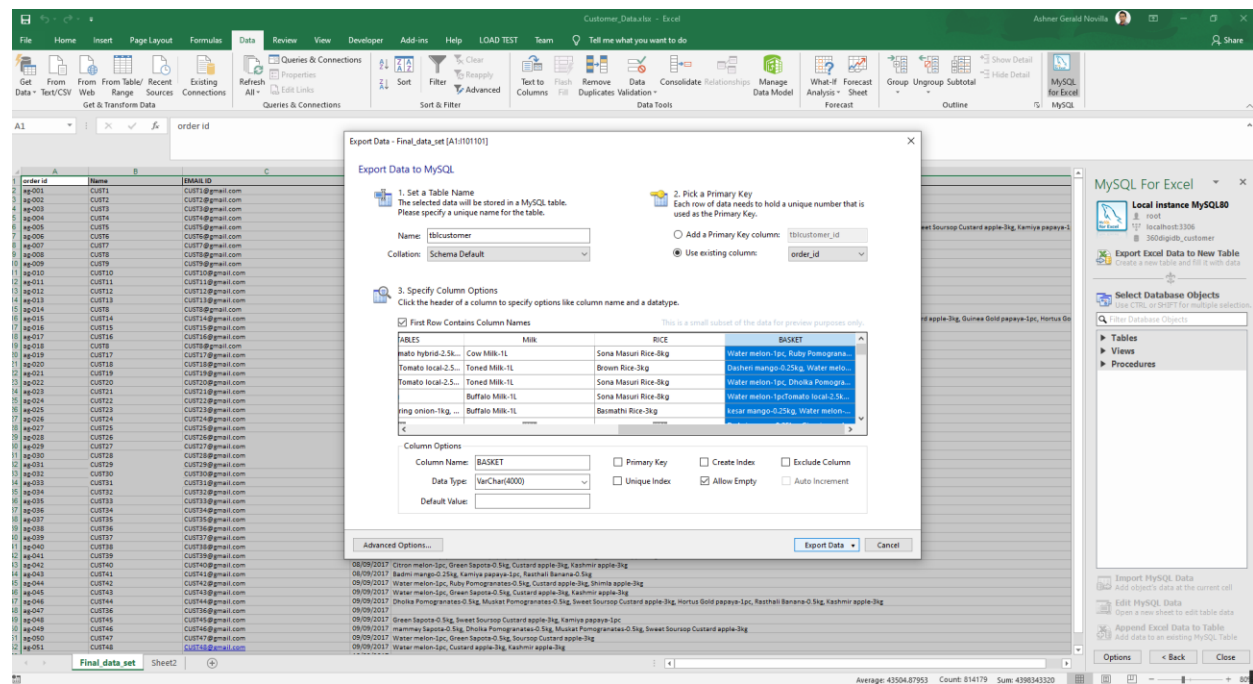
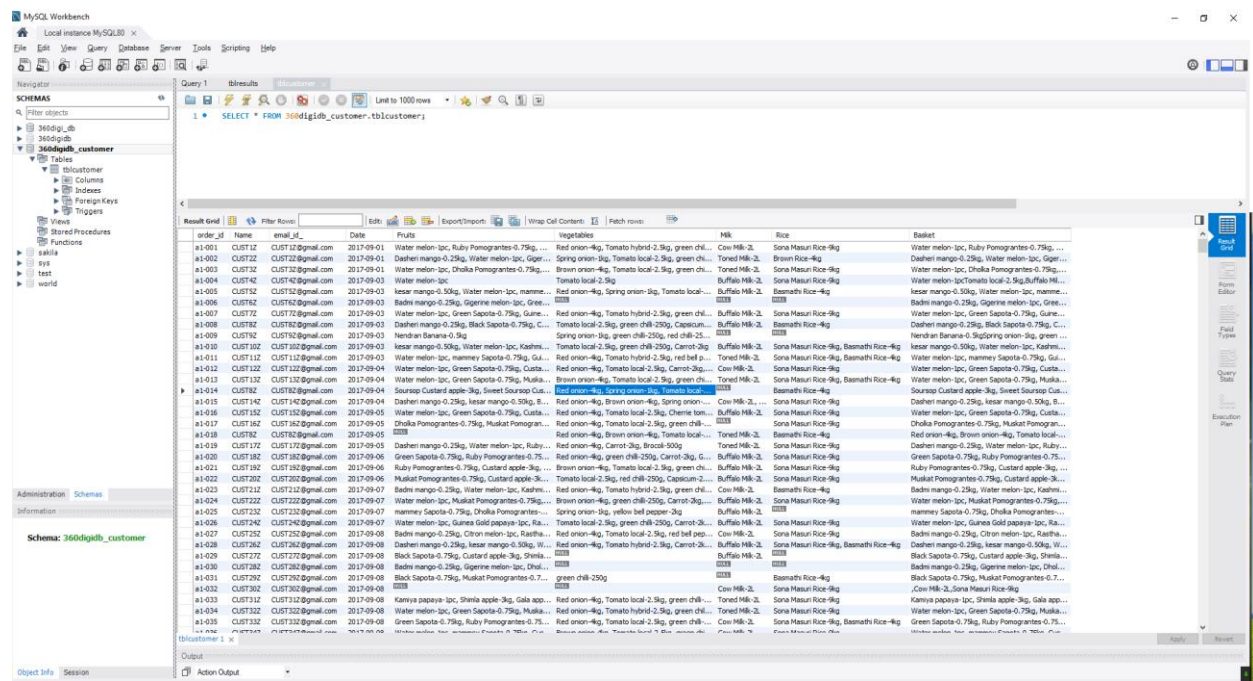


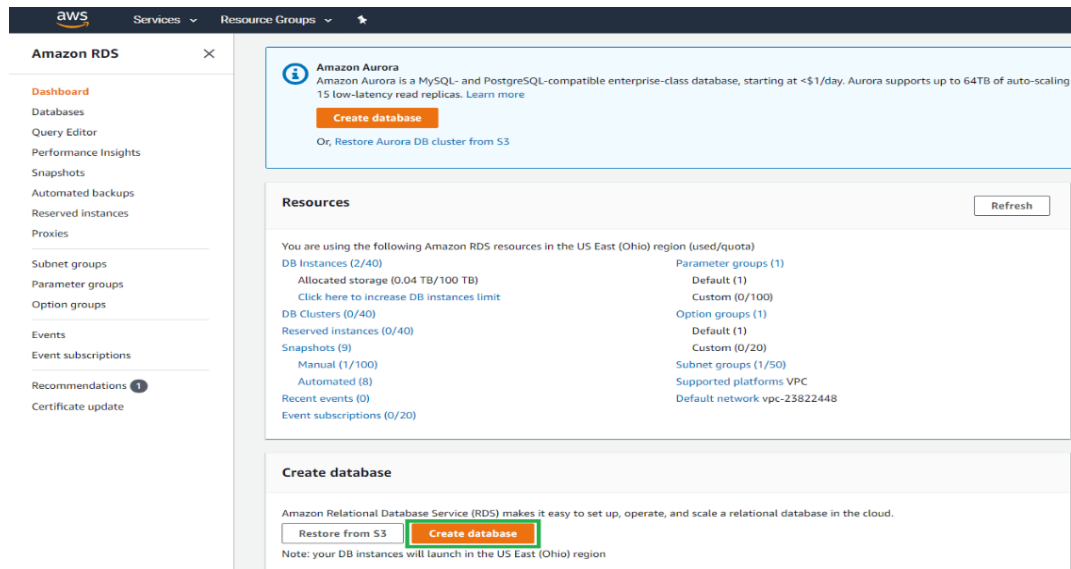
Figure 1.2 Importing table data with MySQL for Excel Successful



2. Steps in Deploying RDS MySQL in Amazon Web Services

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the upper-right corner of the Amazon RDS console, choose the AWS Region in which you want to create the DB instance.
3. In the navigation pane, choose **Databases**.
4. Choose **Create database** and make sure that **Standard Create** is chosen.

Figure 2.1 Initiating the deployment of Database in RDS



5. In **Configuration**, choose **MySQL**.
6. For **DB instance size**, choose **Free tier**.
7. For **DB instance identifier**, enter a name for the DB instance, or leave the default name.
8. For **Master username**, enter a name for the master user, or leave the default name.

The **Create database** page should look similar to the following image.

9. To enter your master password, disable **Auto generate a password**, and then enter the same password in **Master password** and **Confirm password**.
10. Choose **Standard Create** to change one or more settings during database creation and set them.
To change a setting with No in that column, use **Standard Create**. For settings with **Yes** in that column, you can either use **Standard Create** or modify the DB instance after it is created to change the setting.

Figure 2.2 Selecting the Database creation method and the engine option to be used. (In this project the developer will focus on using MySQL, User: administrator, Password: Amazon12#\$)

Choose a database creation method [info](#)

☒ **Standard Create**
You set all of the configuration options, including ones for availability, security, backups, and maintenance.

☐ **Easy Create**
Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type [info](#)

☐ Amazon Aurora

☒ **MySQL**

☐ MariaDB

☐ PostgreSQL

☐ Oracle

☐ Microsoft SQL Server

Edition

☒ **MySQL Community**

Known Issues/Limitations
Review the [Known Issues/Limitations](#) to learn about potential compatibility issues with specific database versions.

Templates

Choose a sample template to meet your use case.

☒ **Production**
Use defaults for high availability and fast, consistent performance.

☐ **Dev/Test**
This instance is intended for development use outside of a production environment.

☐ **Free tier**
Use RDS Free Tier to develop new applications, test existing applications, or gain hands-on experience with Amazon RDS. [info](#)

Settings

DB instance identifier [info](#)

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account in the current AWS Region.

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1 to 60 alphanumeric characters or hyphens (1 to 15 for SQL Server). First character must be a letter. Can't contain two consecutive hyphens. Can't end with a hyphen.

Credentials Settings

Master username [info](#)

Type a login ID for the master user of your DB instance.

1 to 16 alphanumeric characters. First character must be a letter

☐ **Auto generate a password**
Amazon RDS can generate a password for you, or you can specify your own password

Master password [info](#)

Constraints: At least 8 printable ASCII characters. Can't contain any of the following: / (slash), ' (single quote), " (double quote) and @ (at sign).

Confirm password [info](#)

Figure 2.3 Selecting the DB Instance size as the **Standard Class** and having an initial Storage of 100GB SSD with an autoscaling storage up to **1000 GB**. **Multi-AZ Deployment for Active and Standby Mode**.

DB instance size

DB instance class [info](#)

Choose a DB instance class that meets your processing power and memory requirements. The DB instance class options below are limited to those supported by the engine you selected above.

☒ **Standard classes (includes m classes)**

☐ Memory Optimized classes (includes r and x classes)

☐ Burstable classes (includes t classes)

4 vCPUs 16 GiB RAM EBS: 4750 Mbps

☒ **Include previous generation classes**

Storage

Storage type [info](#)

Allocated storage

GiB

Minimum: 100 GiB, Maximum: 65536 GiB

Provisioned IOPS [info](#)

For a workload with 50% writes and 50% reads running on a r3.4xlarge instance, you can realize up to 25,000 IOPS. However, by provisioning more than this limit, you may be able to achieve lower latency and higher throughput. Your actual realized IOPS may vary from the amount you provisioned based on your database workload and instance type. Refer to the [here](#) section to learn more.

Storage autoscaling [info](#)

Provides dynamic scaling support for your database's storage based on your application's needs.

☒ **Enable storage autoscaling**
Enabling this feature will allow the storage to increase once the specified threshold is exceeded.

Maximum storage threshold [info](#)

Charges will apply when your database autoscales to the specified threshold

GiB

Minimum: 101 GiB, Maximum: 65536 GiB

Availability & durability

Multi-AZ deployment [info](#)

☒ **Create a standby instance (recommended for production usage)**
Creates a standby in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

☐ Do not create a standby instance

Connectivity

Virtual private cloud (VPC) [info](#)

VPC that defines the virtual networking environment for this DB instance.

Only VPCs with a corresponding DB subnet group are listed.

After a database is created, you can't change the VPC selection.

Additional connectivity configuration

Subnet group [info](#)

DB subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.

Public access [info](#)

☒ **Yes**
Amazon EC2 instances and devices outside the VPC can connect to your database. Choose one or more VPC security groups that specify which EC2 instances and devices inside the VPC can connect to the database.

☐ **No**
RDS will not assign a public IP address to the database. Only Amazon EC2 instances and devices inside the VPC can connect to your database.

Existing VPC security groups

SG - 360 WDE

default

Database port [info](#)

TCP/IP port that the database will use for application connections.

Figure 2.4 Database authentication using **Password authentication**. The additional configuration set on the image below is configured to cater the ease of **maintenance, back-up, security, monitoring, scaling** and **update** on the client side.

Database authentication

Database authentication options [Info](#)
☒ Password authentication
Authenticates using database passwords.
☐ Password and IAM database authentication
Authenticates using the database password and user credentials through AWS IAM users and roles.
☐ Password and Kerberos authentication (not available for this version)
Choose a directory in which you want to allow authorized users to authenticate with this DB instance using Kerberos Authentication.

▼ Additional configuration
Database options, encryption enabled, backup enabled, backtrack disabled, Performance Insights enabled, Enhanced Monitoring enabled, maintenance, CloudWatch Logs, delete protection enabled

Database options
Initial database name [Info](#)

If you do not specify a database name, Amazon RDS does not create a database.
DB parameter group [Info](#)
default.mysql8.0
Option group [Info](#)
default:mysql-8-0
Backup
Creates a point in time snapshot of your database.
☒ Enable automatic backups
Enabling backups will automatically create backups of your database during a certain time window.

⚠ Please note that automated backups are currently supported for InnoDB storage engine only. If you are using MyISAM, refer to details [here](#).

Backup retention period [Info](#)
Choose the number of days that RDS should retain automatic backups for this instance.
7 days
Backup window [Info](#)
Select the period you want automated backups of the database to be created by Amazon RDS.
☐ Select window
☒ No preference
☒ Copy tags to snapshots
Encryption
☒ Enable Encryption
Choose to encrypt the given instance. Master key IDs and aliases appear in the list after they have been created using the AWS Key Management Service console. [Info](#)
Master key [Info](#)
(default) aws/rds
Account
553219469413
KMS key ID
alias/aws/rds
Performance Insights [Info](#)

ⓘ Enabling Performance Insights will automatically enable the MySQL Community performance schema.
[Learn more](#) [🔗](#).

☒ Enable Performance Insights
Retention period [Info](#)
Default (7 days)
Master key [Info](#)
(default) aws/rds
Account
553219469413
KMS key ID
alias/aws/rds

⚠ You can't change the KMS key after enabling Performance Insights.

Monitoring
☒ Enable Enhanced monitoring
Enabling Enhanced monitoring metrics are useful when you want to see how different processes or threads use the CPU
Granularity
60 seconds
Monitoring Role
default
Clicking "Create database" will authorize RDS to create the IAM role rds-monitoring-role
Log exports
Select the log types to publish to Amazon CloudWatch Logs
☐ Error log
☐ General log
☐ Slow query log
IAM role
The following service-linked role is used for publishing logs to CloudWatch Logs.
RDS service-linked role

ⓘ Ensure that General, Slow Query, and Audit Logs are turned on. Error logs are enabled by default.
[Learn more](#)

Maintenance
Auto minor version upgrade [Info](#)
☒ Enable auto minor version upgrade
Enabling auto minor version upgrade will automatically upgrade to new minor versions as they are released. The automatic upgrades occur during the maintenance window for the database.
Maintenance window [Info](#)
Select the period you want pending modifications or maintenance applied to the database by Amazon RDS.
☐ Select window
☒ No preference
Deletion protection
☒ Enable deletion protection
Protects the database from being deleted accidentally. While this option is enabled, you can't delete the database.

Figure 2.5 Database estimated monthly costs that the client will need to keep in mind. (link: <https://calculator.s3.amazonaws.com/index.html> for future calculation or resizing of the environment)

Estimated monthly costs	
DB instance	249.66 USD
Storage	25.00 USD
Multi-AZ standby instance	249.66 USD
Provisioned IOPS	100.00 USD
Total	624.32 USD
This billing estimate is based on on-demand usage as described in Amazon RDS Pricing . Estimate does not include costs for backup storage, IOs (if applicable), or data transfer.	
Estimate your monthly costs for the DB Instance using the AWS Simple Monthly Calculator .	

3. Connecting to a Database on a DB Instance Running the MySQL Database Engine

- After Amazon RDS provisions your DB instance, you can use any standard SQL client application to connect to a database on the DB instance. In this example, you connect to a database on a MySQL DB instance using MySQL monitor commands.
- Find the endpoint (DNS name) and port number for your DB instance.
 - Open the RDS console and then choose Databases to display a list of your DB instances.
 - Choose the MySQL DB instance name to display its details.
 - On the Connectivity & security tab, copy the endpoint. Also, note the port number. You need both the endpoint and the port number to connect to the DB instance.

Figure 3.1 Database Endpoint (*this endpoint is been done using free tier mode for the developer to avoid unneeded cost of deployment*)

aswdbcloudinstance

Summary

DB identifier

aswdbcloudinstance

CPU

2.03%

Info

Available

Class

db.t2.micro

Role

Instance

Current activity

0 Connections

Engine

MySQL Community

Region & AZ

us-east-2b

Connectivity & security

Monitoring

Logs & events

Configuration

Maintenance & backups

Tags

Connectivity & security

Endpoint & port

Endpoint

aswdbcloudinstance.cdldqzpuco.us-east-2.rds.amazonaws.com

Port

3306

Networking

Availability zone

us-east-2b

VPC

vpc-23822448

Subnet group

default-vpc-23822448

Subnets

subnet-09177845
subnet-ab9a9ed1
subnet-e610de8d

Security

VPC security groups

SG - 360 WDE (sg-07f3fdae858789d0
(active)
default: sg-e6d2c99c
(active)

Public accessibility

Yes

Certificate authority

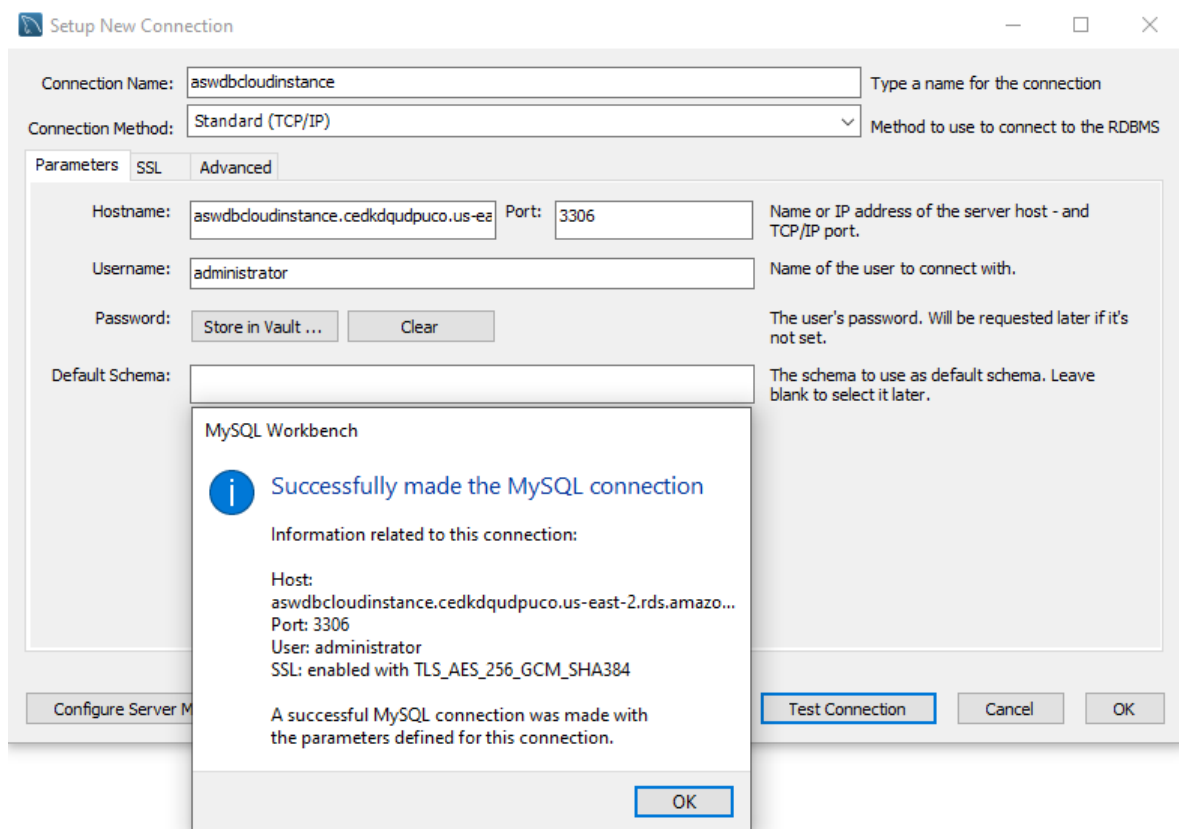
rd5-ca-2019

Certificate authority date

Aug 22nd, 2024

3. Download a SQL client that you can use to connect to the DB instance. You can connect to an Amazon RDS MySQL DB instance by using tools like the MySQL Workbench.
4. Connect to a database on a MySQL DB instance. For example, enter the following command at a command prompt on a client computer to connect to a database on a MySQL DB instance using the MySQL client. Substitute the DNS name for your DB instance for *<endpoint>*, the master user name you used for *<mymasteruser>*, and provide the master password you used when prompted for a password.

Figure 3.2 Local MySQL Workbench successfully made a connect to MySQL of RDS in AWS

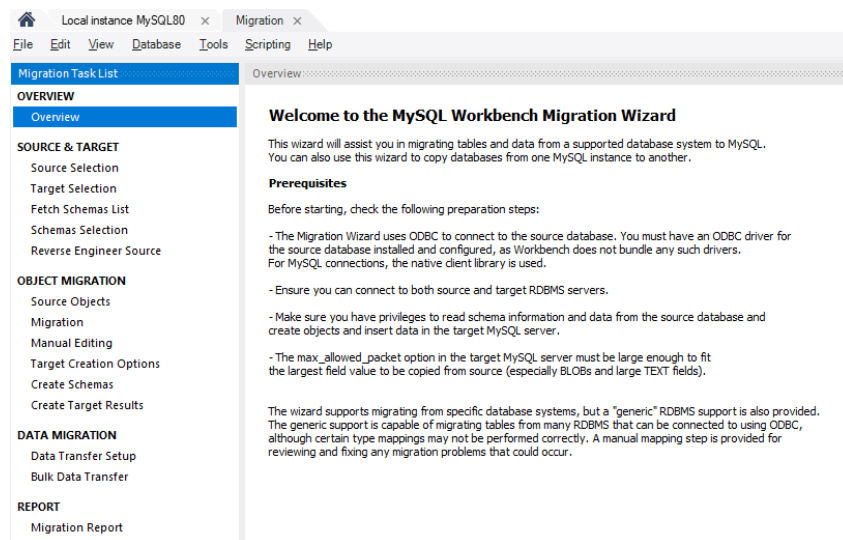


4. Migrating an On-Premises Database to Amazon RDS MySQL

The MySQL Workbench Migration Wizard allows you to easily and quickly migrate databases from Microsoft SQL Server, PostgreSQL, Sybase ASE and SQL anywhere, SQLite, and most ODBC-capable RDBMSs to MySQL. In addition, you can use the module to create MySQL to MySQL database copies that can be used for tasks such as copying a database across servers or migrating data across different versions of MySQL.

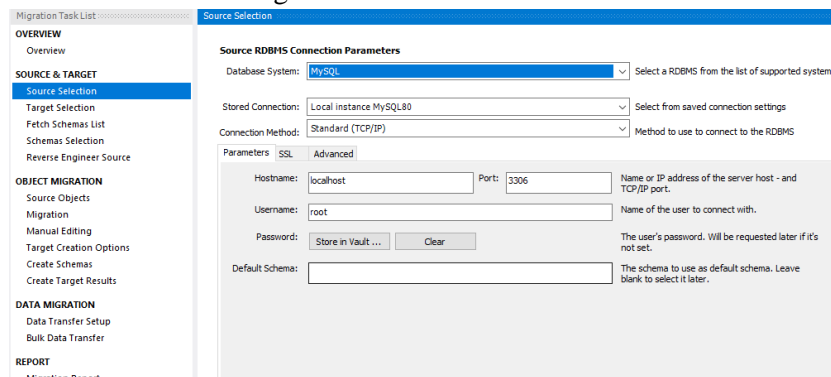
1. Check if the user has an appropriate access to a running SQL Server instance of the database you want to migrate.
2. Check if the user has an appropriate access to a running MySQL Server instance.
3. Open MySQL Workbench and start the Migration Wizard. From the main MySQL Workbench screen, you can start the Migration Wizard by clicking on the Database Migration launcher in the Workbench Central panel or through **Database** → **Migrate** in the main menu.

Figure 4.1 Migration Wizard



4. Set up the parameters to connect to your source database by **clicking start**. In this page you need to provide the information about the RDBMS you are migrating.
5. If you open the **Database System** combo box, you'll find a list of the supported RDBMSs. Select MySQL Server from the list. Just below it there's another combo box named **Stored Connection**. It will list saved connection settings for that RDBMS. You can save connections by marking the checkbox at the bottom of the page and giving a name.

Figure 4.2 Source Connection Settings



- Set up the parameters to connect to your target Database. Click on the **Next** button to move to the **Target Selection** page. Once there, set the parameters to connect to your MySQL Server instance. When you are done click on the **Test Connection button** and verify that you can successfully connect to it.

Figure 4.3 Target Connection Settings

The screenshot shows the 'Target Selection' page of a migration tool. On the left is a sidebar with a 'Migration Task List' containing sections: OVERVIEW (Overview), SOURCE & TARGET (Source Selection, Target Selection, Fetch Schemas List, Schemas Selection, Reverse Engineer Source), OBJECT MIGRATION (Source Objects, Migration, Manual Editing, Target Creation Options, Create Schemas, Create Target Results), DATA MIGRATION (Data Transfer Setup, Bulk Data Transfer), and REPORT (Migration Report). The 'Target Selection' item is highlighted. The main panel is titled 'Target RDBMS Connection Parameters' and includes a 'Stored Connection' dropdown (set to 'aswdbcloudinstance'), a 'Connection Method' dropdown (set to 'Standard (TCP/IP)'), and tabs for 'Parameters', 'SSL', and 'Advanced'. The 'Parameters' tab is active, showing fields for 'Hostname' (aswdbcloudinstance.cedkqdpuco.us-e), 'Port' (3306), 'Username' (administrator), 'Password' (with 'Store in Vault' and 'Clear' buttons), and 'Default Schema'. To the right of these fields are explanatory text blocks: 'Name or IP address of the server host - and TCP/IP port.', 'Name of the user to connect with.', 'The user's password. Will be requested later if it's not set.', and 'The schema to use as default schema. Leave blank to select it later.'

Figure 4.4 Fetch Schemas List

The screenshot shows the 'Fetch Schema List' page. The sidebar is identical to the previous figure, with 'Fetch Schemas List' highlighted under the 'SOURCE & TARGET' section. The main panel has a title bar 'Fetch Schema List' and contains the following text: 'The following tasks will now be performed. Please monitor the execution.', 'The names of available schemas will be retrieved from the source RDBMS. The account used for the connection will need to have appropriate privileges for listing and reading the schemas you want to migrate. Target RDBMS connection settings will also be checked for validity.', a list of three checked tasks: 'Connect to source DBMS', 'Check target DBMS connection', and 'Retrieve schema list from source', and a final instruction: 'Finished performing tasks. Click [Next >] to continue.'

7. Select the Schema to Migrate Click on the Next button to move to the next page. The Migration Wizard will connect to SQL Server to fetch a list of the catalogs and schemas.

Figure 4.4 Schemas Selection (In this project the user will be importing the data from 360digidb_customer schema that contains the data given by the PM)

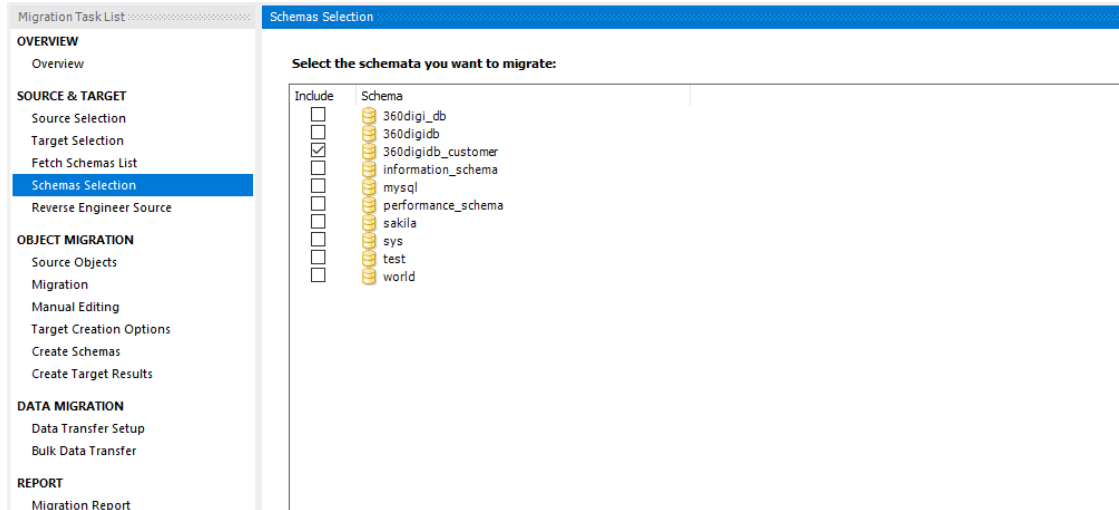
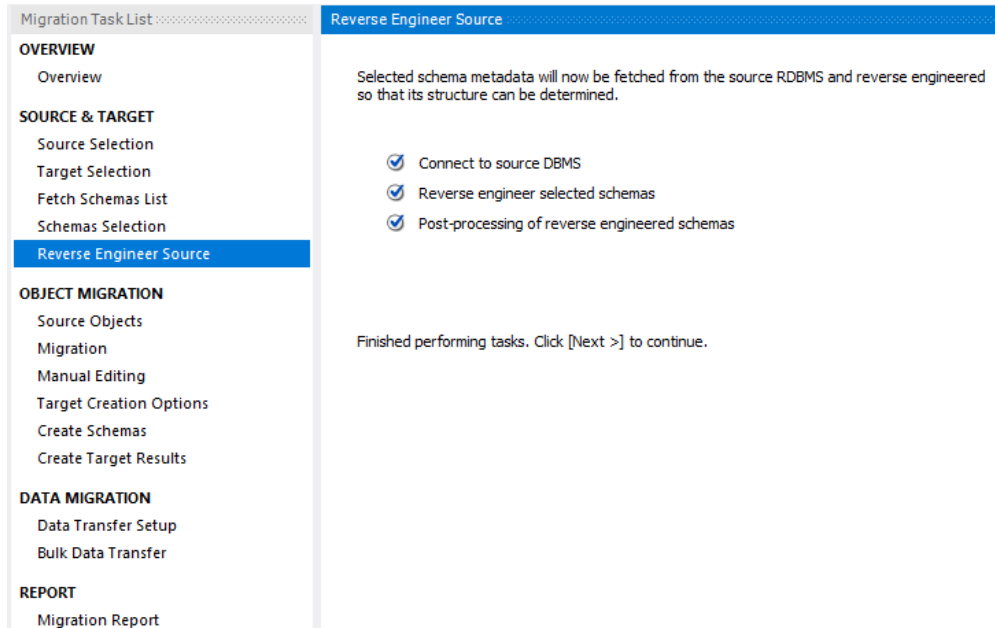
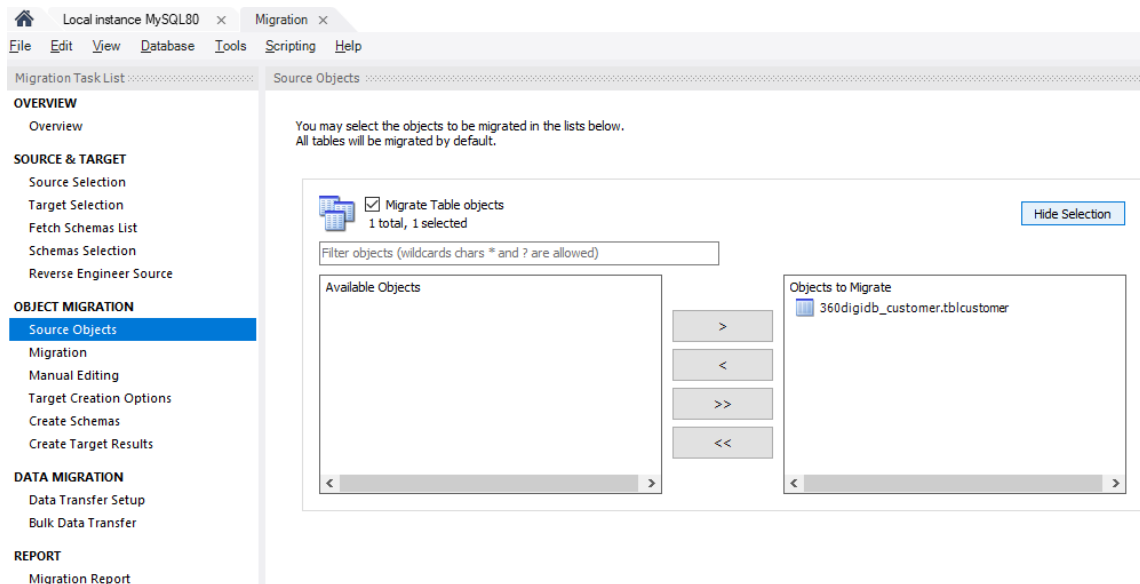


Figure 4.5 Reverse Engineer Source



8. Select the Objects to Migrate Move to the next page using the Next button. You should see the reverse engineering of the selected schema in progress. At this point the Migration Wizard is retrieving relevant information about the involved database objects (table names, table columns, primary and foreign keys, indices, triggers, views, etc.).

9. If you click on the **Show Selection** button you will be given the opportunity to select exactly what you want to migrate as shown here:



10. Review the Proposed Migration. Move to the next page. You will see the progress of the migration there. At this point the Migration Wizard is converting the objects you selected into their equivalent objects in MySQL and creating the MySQL code needed to create them in the target server.

Figure 4.5 Migration

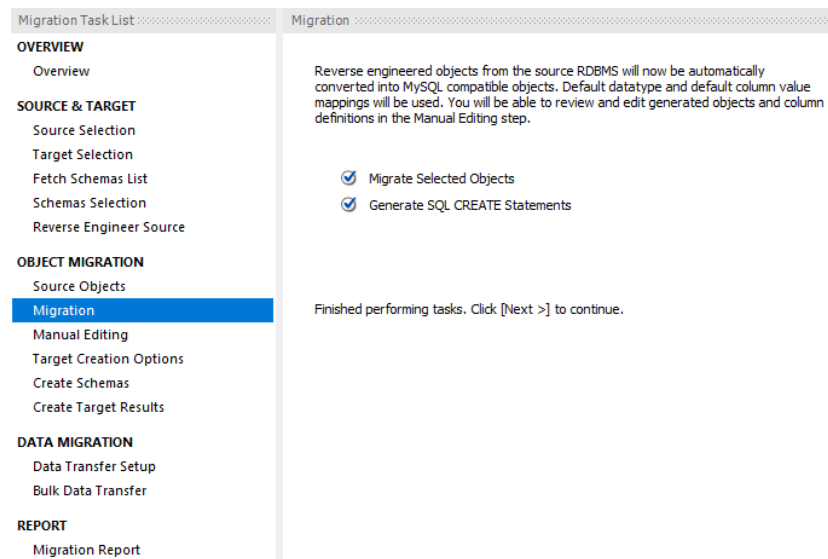


Figure 4.6 Manual Editing

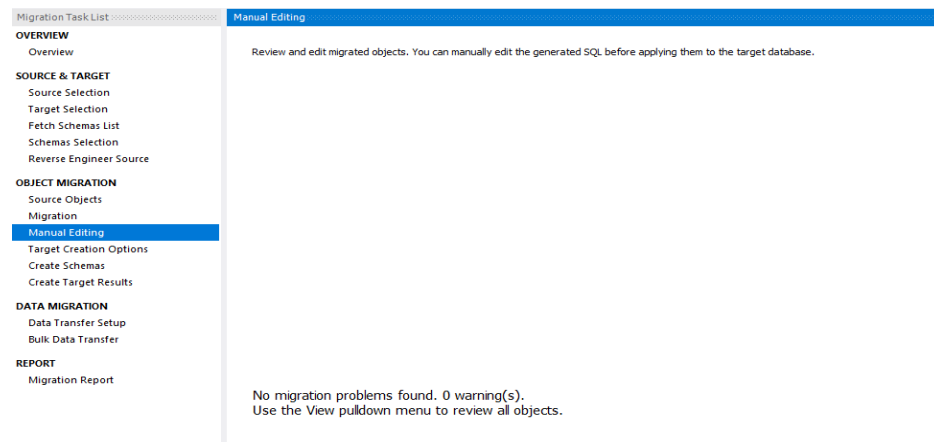


Figure 4.7 Target Creation Options

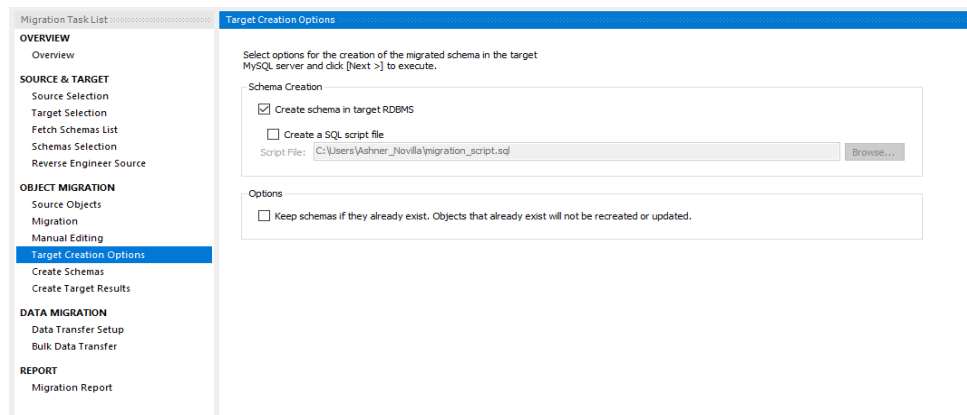


Figure 4.8 Create Schemas

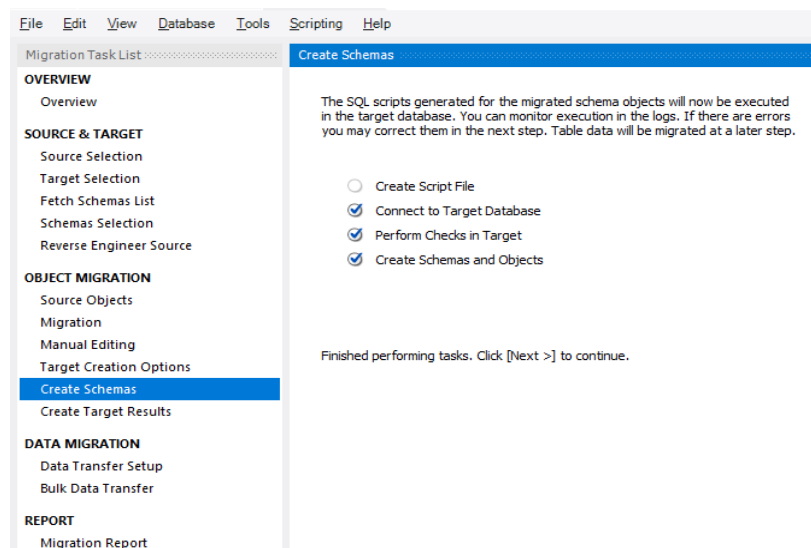
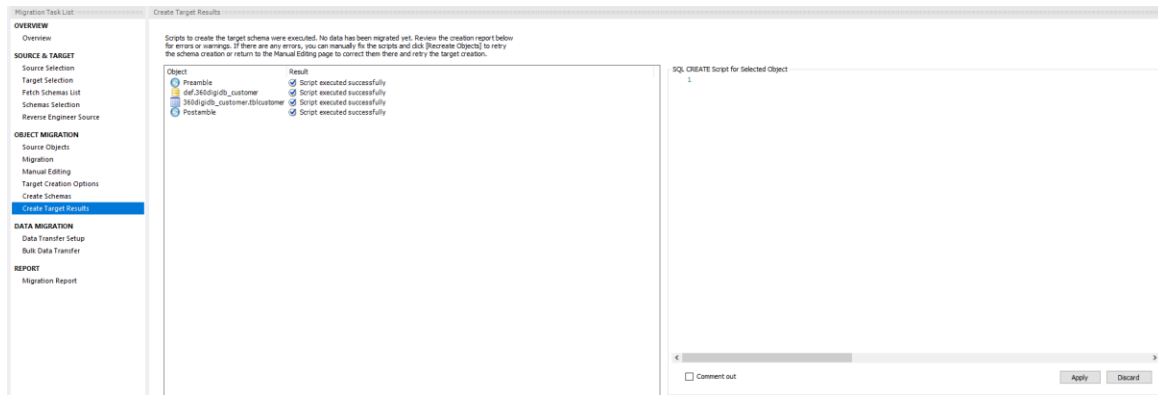


Figure 4.96 Create Target Results



11. This step in the Migration Wizard involve the transference of data from the source SQL Server database into your newly created MySQL database. The Data Transfer Setup page allows you to configure this process.

Figure 4.97 Data Transfer Setup

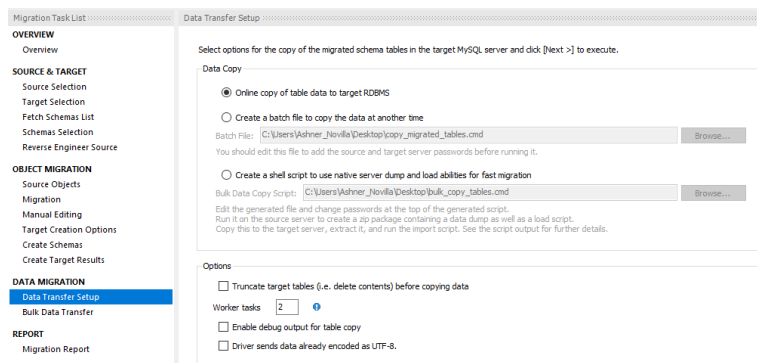
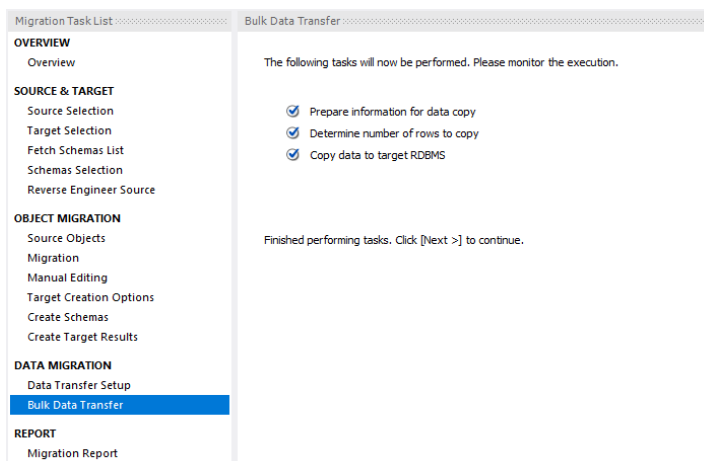


Figure 4.98 Bulk Data Transfer



12. Once it finishes, move to the next page. A report page will be displayed with a summary of the entire process:

Figure 4.98 Migration Report and Verification of Migration Result on the MySQL RDS in AWS

The screenshot displays the MySQL Workbench Migration Wizard interface, showing the 'MigrationReport' tab. The report is titled 'MySQL Workbench Migration Wizard Report' and includes the following sections:

- OVERVIEW**
 - Overview
- SOURCE & TARGET**
 - Source Selection
 - Target Selection
 - Fetch Schemas List
 - Schemas Selection
 - Reverse Engineer Source
- OBJECT MIGRATION**
 - Source Objects
 - Migration
 - Manual Editing
 - Target Creation Options
 - Create Schemas
 - Create Target Results
- DATA MIGRATION**
 - Data Transfer Setup
 - Bulk Data Transfer
- REPORT**
 - Migration Report

The report content includes:

- 1. Migration**
 - 1. Summary**
 - Number of migrated schemas: 1
 - 1. 360digdb_customer
 - Source Schema: 360digdb_customer
 - Tables: 1
 - Triggers: 0
 - Views: 0
 - Stored Procedures: 0
 - Functions: 0
 - 2. Migration Issues**
 - 3. Object Creation Issues**
 - 4. Migration Details**
 - 4.1. Table 360digdb_customer.tblcustomer (tblcustomer)
 - Columns:
 - order_id VARCHAR(12)
 - Name VARCHAR(12)
 - email_id VARCHAR(25) NULL
 - Date DATE NULL
 - Fruits VARCHAR(4000) NULL
 - Vegetables VARCHAR(4000) NULL
 - Milk VARCHAR(4000) NULL
 - Rice VARCHAR(4000) NULL
 - Basket TEXT NULL
 - Foreign Keys:
 - Indices:
 - PRIMARY (order_id)
- II. Data Copy**
 - '360digdb_customer'. 'tblcustomer'
 - Succeeded : copied 101100 of 101100 rows from '360digdb_customer'. 'tblcustomer'

The bottom section of the screenshot shows the 'Administration - Server Status' window with a query executed: `SELECT * FROM 360digdb_customer.tblcustomer;`. The query results are displayed in a table with columns: order_id, Name, email_id, Date, Fruits, Vegetables, Milk, and Rice. The table contains 10 rows of data, showing various food items and their details.